



---

# *Computer Networks*

## *Assignment 2: Router*

### *Tutorial*

*Tutorial: Oct 12<sup>th</sup>, 2022*

*Due Date: 5pm Nov 1<sup>st</sup>, 2022*


---

# **PART 1: ASSIGNMENT DESCRIPTION**

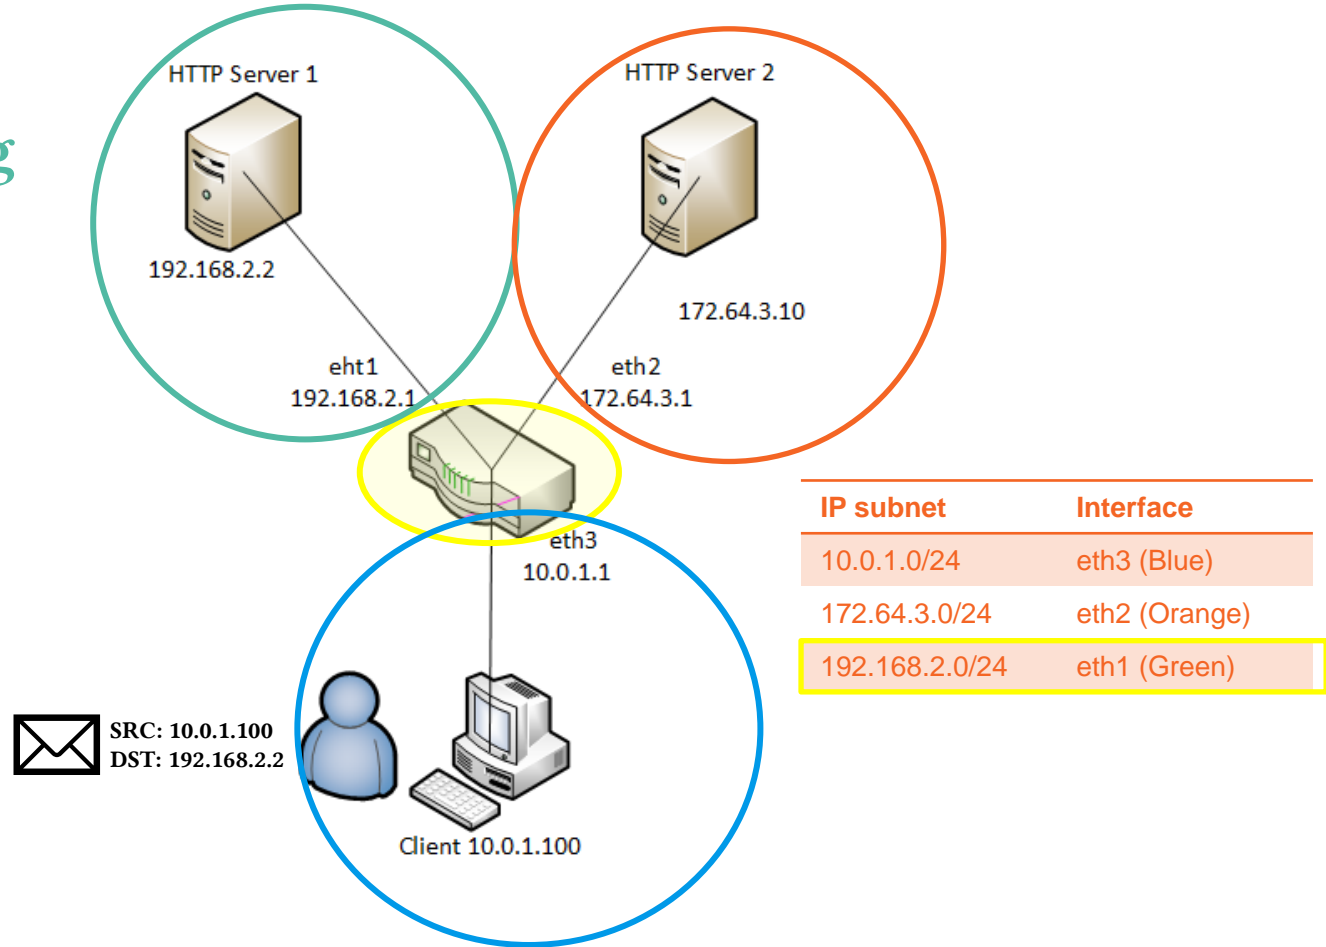
## **Implementing a simple router**

# Overview

- You will build a **simplified router** which will route packets from an emulated client to two emulated servers.
- Emulation environment: [Mininet](#)
- Three main parts to the assignment:
  - Preparing the emulation environment
  - Implementing IP forwarding
  - Handling address resolution (ARP)



# Overview: IP Forwarding



# Overview: ARP Handling



192.168.2.2 IS AT  
ab:cd:ef:ab:cd:ef

HTTP Server 1



192.168.2.2

eth1

192.168.2.1

HTTP Server 2



172.64.3.10


eth2

172.64.3.1

I don't know the  
MAC address for  
192.168.2.2!




WHO HAS  
192.168.2.2?  
DST: 192.168.2.2




eth3

10.0.1.1

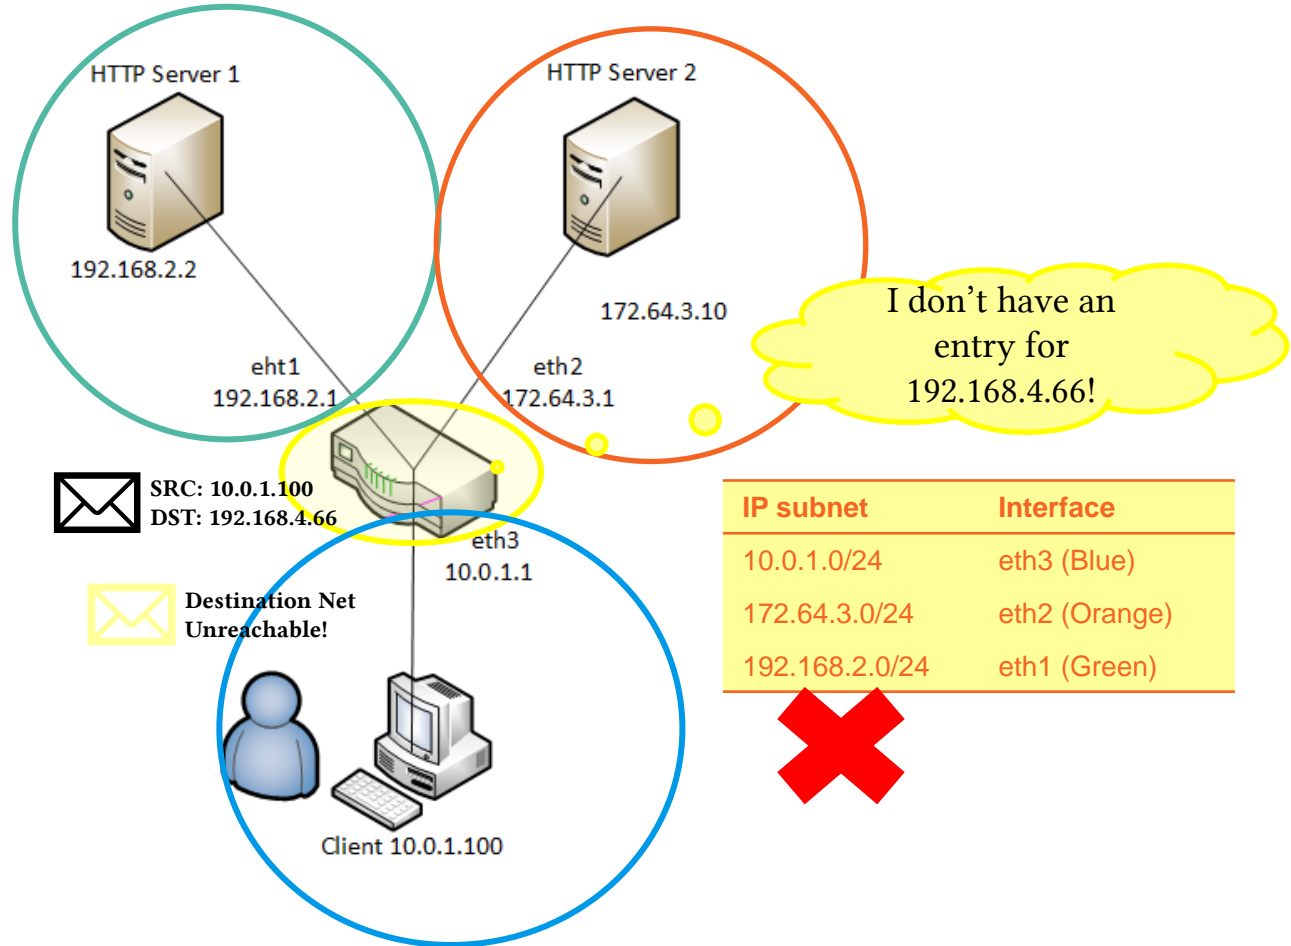


Client 10.0.1.100

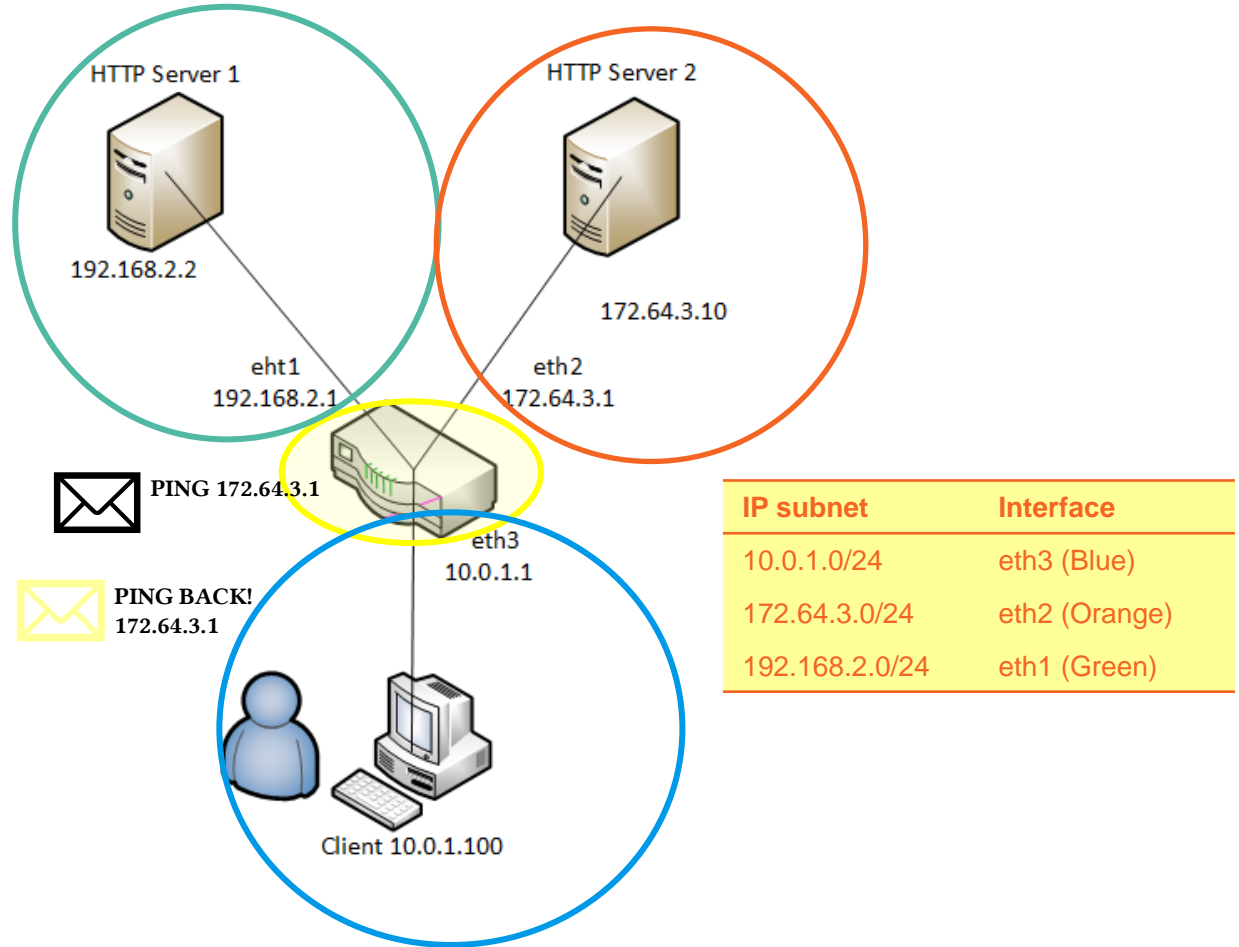
| IP subnet      | Interface     |
|----------------|---------------|
| 10.0.1.0/24    | eth3 (Blue)   |
| 172.64.3.0/24  | eth2 (Orange) |
| 192.168.2.0/24 | eth1 (Green)  |



# Overview: ICMP handling



# Overview: ICMP handling



# Raw Ethernet Frames

- Your router will receive and send **raw Ethernet** frames.



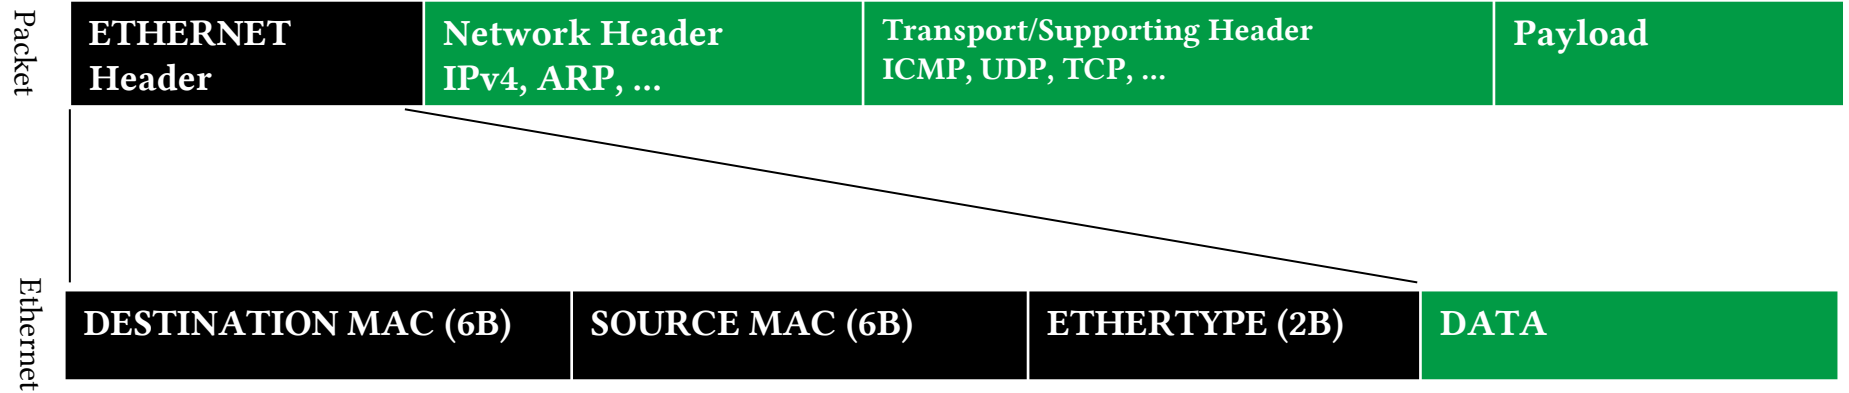
What the router software receives



# Ethernet

- Your router will receive and send **raw Ethernet** frames.
- Ethernet frames have headers which contain source and destination MAC addresses.
  - To forward a packet one hop, we must **change** the destination MAC address of the forwarded packet to the MAC address of the next hop's incoming interface.
- Your router must **process** packets & **forward** them to the correct interface.
- When your router receives a packet
  - First, **identify** the packet (IP, ARP)
  - Then, **extract** appropriate header

# Ethernet Header

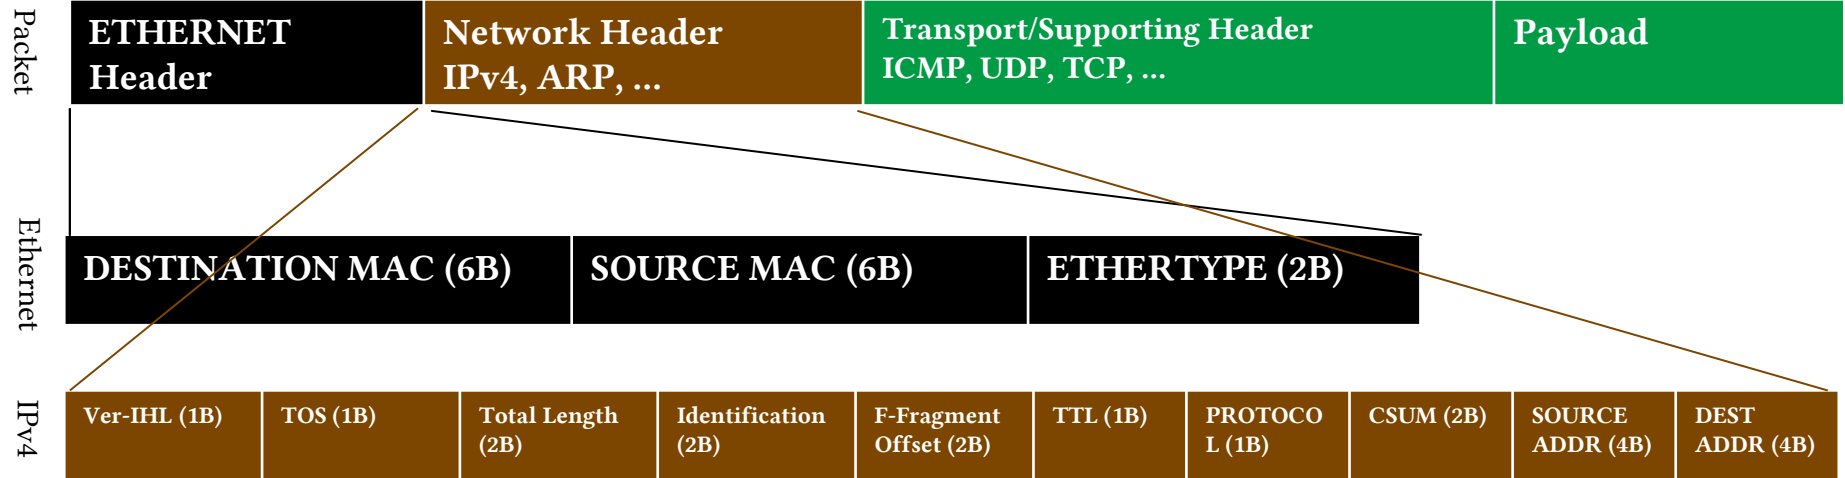


# Internet Protocol (IPv4)

If your router receives an IPv4 packet:

1. Perform basic checks
  - a. Verify its checksum
  - b. Check that it meets minimum length
2. Determine if it is destined for one of our interfaces:
  - a. If yes, check that it is an ICMP echo request (type 8)
    - i. If yes, send **ICMP echo reply (type 0)** back
    - ii. Else, drop it. If not even an ICMP packet, send **ICMP port unreachable (type 3, code 3)**
  - b. If not destined for our interfaces, IPv4 packet **must be forwarded**.

# IPv4 Header



# IP Forwarding

If packet must be forwarded:

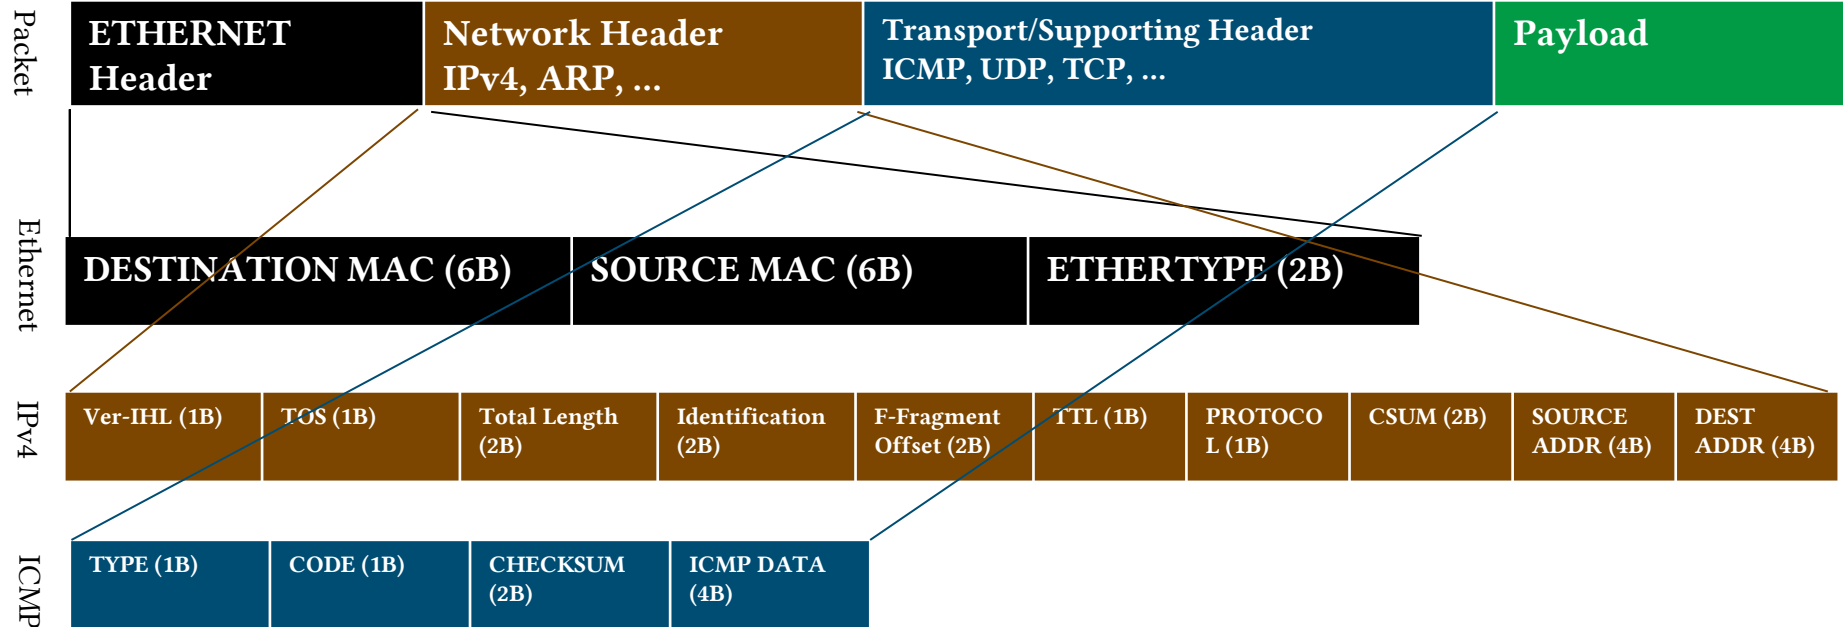
1. Decrement TTL by 1, recompute packet checksum
  - a. If TTL is expired, send **ICMP time exceeded (type 11)**
2. Find match for the destination IP address in the routing table
  - a. If no match, send **ICMP destination net unreachable (type 3, code 0)**
  - b. If match, check ARP cache for next-hop MAC address corresponding to next-hop IP
    - i. If ARP entry is found, **send IP packet to its next hop**
    - ii. If no ARP entry is found:
      1. Send ARP request for next-hop IP
      2. Put **packet onto a queue** as it waits for an ARP reply.

# ICMP

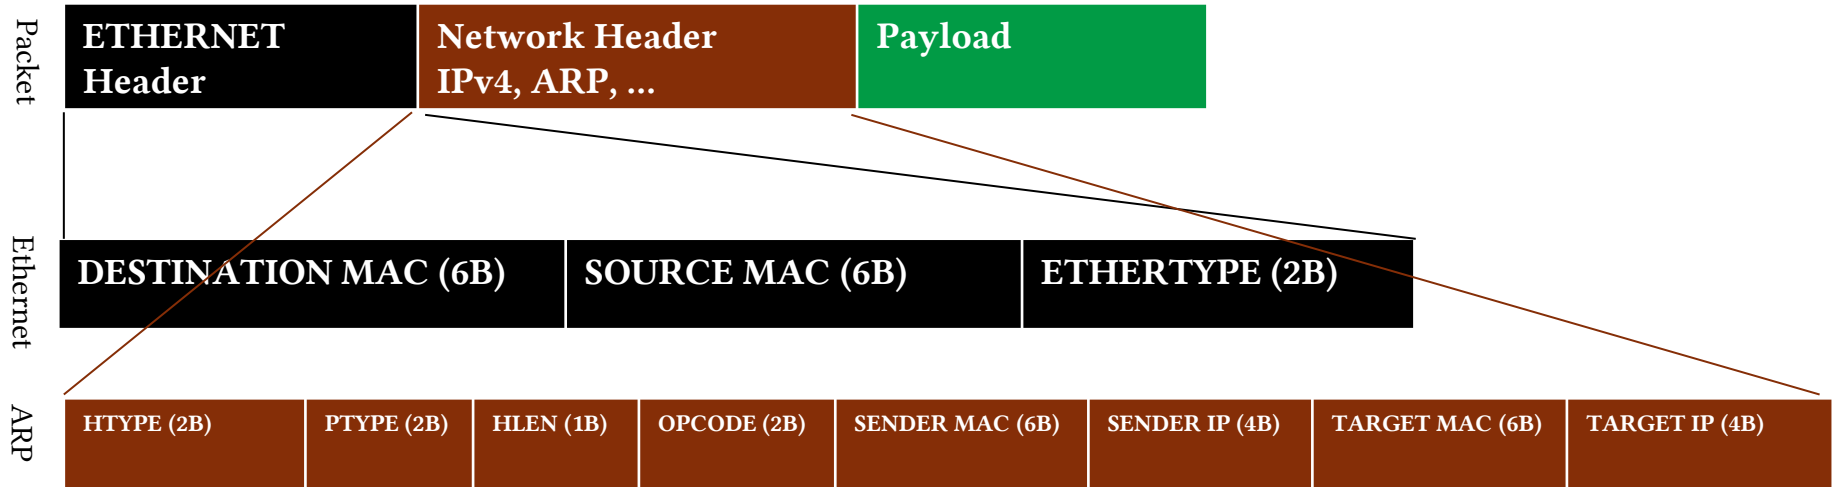
Use ICMP messages to send control information to host:

- **Echo reply (type 0)**
  - Send in response to echo request (type 8). Used for PING.
- **Destination net unreachable (type 3, code 0)**
  - Sent if there is a non-existent route to the destination IP.
- **Destination host unreachable (type 3, code 1)**
  - Sent if five ARP requests were sent to the next-hop IP without a response.
- **Port unreachable (type 3, code 3)**
  - Sent if an IP packet containing a UDP or TCP payload is sent to one of the router's interfaces.
- **Time exceeded (type 11, code 0)**
  - Sent if an IP packet is discarded because the TTL field is 0.

# ICMP Header



# ARP Header





# Address Resolution Protocol (ARP)

- ARP determines the next-hop MAC that corresponds to the next-hop IP
- Your router needs to generate and process ARP requests/replies.
  - You are required to cache ARP replies that your router receives
  - We handle the cache timeout for you
- During the packet forwarding process, the router should first check the ARP cache for the MAC address before sending an ARP request

# Address Resolution Protocol (ARP)

If your router receives an ARP packet:

- Check if it's an ARP request or ARP reply:
  - If ARP request:
    - Check if its target IP address is one of our router's IP addresses.
      - If yes, **send ARP reply**
      - Else, drop it
  - If ARP reply:
    - **Cache the ARP reply** as long as the target IP address is one of your router's IP addresses
    - If there is an ARP request waiting on the reply
      - **Send IP packets** waiting on this ARP reply
      - Remove corresponding ARP request from queue

# END OF PART 1: Assignment Description

Q&A

# **PART 2: TESTING ENVIRONMENT**

## **Deploying Mininet using Vagrant & Getting started**

# Mininet VM Setup

- Checkout the appropriate branch on [Github](#)
- Custom Vagrant VM setup (in ./Assignment2/ dir) → **must use this one!**
- Mininet emulator is *already set up* for you, so it should be painless :)
- Test connectivity using the reference binary provided.
  - **Do this BEFORE writing any code.**
  - **Do this early** so we can help resolve any setup issues.
  - If there are any connectivity issues, you will not be able to test your implementation.
- Step-by-step instructions are in the README

# Getting Started: sr\_router.c

- In `sr_router.c`:

```
void sr_handlepacket(struct sr_instance* sr, uint8_t *  
packet, unsigned int len, char* interface)
```

- This method is called every time the router receives a raw Ethernet frame.
- You will have to implement the logic to process the packets and send them to the correct interface.
  - **Tip:** First check if it is an IP or an ARP packet using `ethertype()` method from `sr_utils.c/h`, then process accordingly
  - Use helper methods!

# Helper functions: sending a packet

- In `sr_vns_comm.c`:

```
int sr_send_packet(struct sr_instance* sr, uint8_t* buf,  
unsigned int len, const char* iface)
```

- Sends a packet of length `len` to the network out of `iface`
- Your `sr_router.c` should call this method
- You do not need to modify this method

# Getting Started: sr\_arpcache.c

- In `sr_arpcache.c`:

```
void sr_arpcache_sweepreqs(struct sr_instance *sr)
```

- Called once per second
- You should add code to iterate through the ARP request queue and re-send outstanding requests
- Fill in the helper method called e.g. `handle_arpreq()` to track of how many times the request has been sent
  - If sent  $\geq 5$  times and no response, send **ICMP destination host unreachable (type 3, code 1)**
- **Tip:** there is helpful pseudocode for this in `sr_arpcache.h`



# Summary

- Code to be written in `sr_router.c`:
  - Handle IP packet
    - Send ICMP echo replies
    - Send (forward) IP packets
    - Queue ARP requests
  - Handle ARP packet
    - Send ARP replies in response to requests
    - Cache incoming ARP reply
    - Send IP packets that are waiting on incoming ARP reply
    - Remove corresponding ARP request from queue
  - Send ICMP messages based on certain conditions
- Code to be written in `sr_arpcache.c`:
  - Iterate through ARP request queue
  - Send and re-send ARP requests
  - Check if we need to re-send or destroy ARP request
- There are some helpful methods to look up interfaces in `sr_if.c`.
- There are some generally helpful methods in `sr_utils.c`.

**These are just my recommendations for where to place different functionality. You may organize the functionality across the code base as you see fit.**

# Final Thoughts

- Please start early.
  - I repeat, **start early.**
- Come to office hours
- Ask questions on Canvas

