

Assignment 3- Transport Layer

11/4/2022

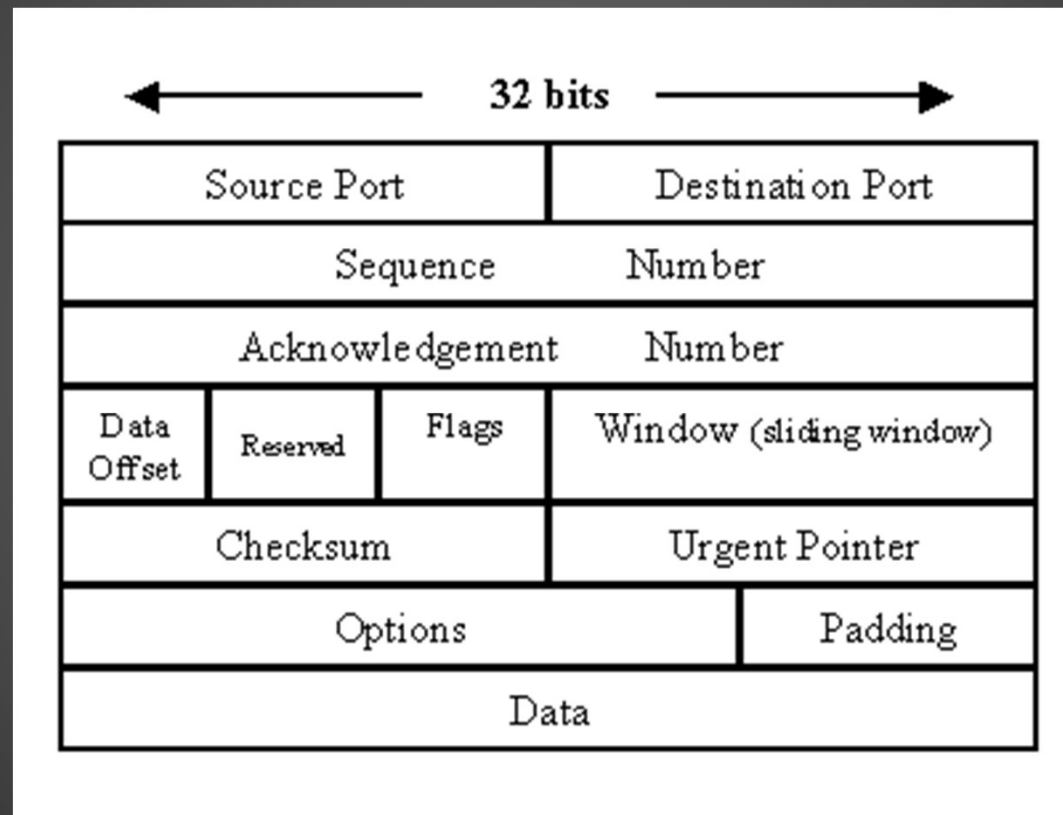
Overview

- ▶ Implement simplified TCP known as STCP
 - ▶ Runs on custom socket interface called MYSOCK
 - ▶ Assume reliable network layer where in-order delivery is guaranteed:
 - ▶ No dropped packets or retransmissions
 - ▶ No reordering
 - ▶ No timeouts
 - ▶ Does not include congestion control
 - ▶ Receives data from application and sends to networks
 - ▶ Receives data from network and sends to application

STCP Protocol

- ▶ Connection-Oriented:
 - ▶ Handshake establishing connection parameters must be performed before beginning sequential data transfer
- ▶ Data Treated as a Stream:
 - ▶ Must read all data sent by one peer, break the stream into packets and reassemble the stream on the receiving peer's side
- ▶ Full-Duplex:
 - ▶ Each connection is a pair of byte streams, one for each connection side
 - ▶ Sending on connection goes both ways
 - ▶ Active- Sending/Connection Initiator
 - ▶ Passive- Receiving

TCP Packet



TCP Packet Structure

- ▶ Fields used by STCP Implementation:
 - ▶ th_seq
 - ▶ Sequence number of first byte in payload
 - ▶ th_ack
 - ▶ The sequence number being acknowledged if this is an ACK packet
 - ▶ th_off
 - ▶ The offset within the packet at which data begins
 - ▶ No optional data in this assignment
 - ▶ th_flags
 - ▶ Any flags OR'd together
 - ▶ th_win
 - ▶ Size of advertised receiver window in bytes

```
typedef uint32_t tcp_seq;

struct tcphdr {
    uint16_t th_sport;           /* source port */
    uint16_t th_dport;           /* destination port */
    tcp_seq th_seq;              /* sequence number */
    tcp_seq th_ack;              /* acknowledgment number */
#ifdef _BIT_FIELDS_LTOH
    u_int   th_x2:4,             /* (unused) */
            th_off:4;           /* data offset */
#else
    u_int   th_off:4,            /* data offset */
            th_x2:4;            /* (unused) */
#endif
    uint8_t th_flags;
#define TH_FIN  0x01
#define TH_SYN  0x02
#define TH_RST  0x04
#define TH_PUSH 0x08
#define TH_ACK  0x10
#define TH_URG  0x20
    uint16_t th_win;             /* window */
    uint16_t th_sum;             /* checksum */
    uint16_t th_urp;             /* urgent pointer */
    /* options follow */
};
```

Sequence Number

- ▶ Assigns by numbering bytes

- ▶ Rules:

- ▶ Sequentially number
 - ▶ Sequence Number = $n + \text{Initial Sequence Number}$
 - ▶ n is the number of bytes in the bytestream
- ▶ Sequence initialized to random number within 0-255 (inclusive)
- ▶ Should be set for every packet

- ▶ SYN & FIN flags

- ▶ Also set with the sequence numbers
- ▶ Associated with 1 byte each of the sequence space
- ▶ SYN- synchronizes sequence numbers of peers
- ▶ FIN- indicates the end of the communication

- ▶ Packet Data:

- ▶ Maximum STCP packet payload = 536 bytes
 - ▶ Stored as STCP_MSS
- ▶ Send data as soon as available and if its within the effective window

```
▼ Transmission Control Protocol, Src Port: 80
  Source Port: 80
  Destination Port: 60862
  [Stream index: 16]
  [TCP Segment Len: 1025]
  Sequence number: 403851 (relative sequ
  [Next sequence number: 404876 (relativ
  Acknowledgment number: 434394 (relativ
  Header Length: 20 bytes
  ▶ Flags: 0x018 (PSH, ACK)
```


ACK

- ▶ Ensures reliable delivery
- ▶ After receiving packet, the receiver replies with an ACK packet to indicate that the packet was received
- ▶ Refer to the sequence number of the next expected byte of data
 - ▶ Example:
 - ▶ Sender sent bytes 512-1023
 - ▶ ACK reply = 1024 + Initial Sequence Number
- ▶ ACK packets contain 0 bytes of payload
- ▶ Rules:
 - ▶ Send ACK as soon as data received
 - ▶ Unlike TCP which delays sending the ACK
 - ▶ If packet has duplicate data, send a new ACK for the next expected sequence number

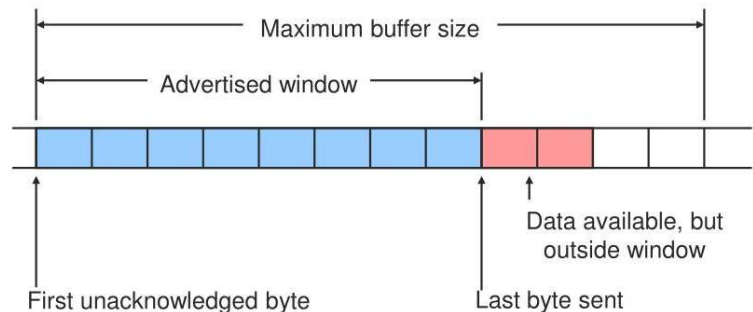
Sliding Window

- ▶ Tracks how many unacknowledged “in-flight” packets there are
 - ▶ Window size “slides” by incrementing when an ACK is received
- ▶ Receiver Window:
 - ▶ Prevents sender from overwhelming receiver
 - ▶ Sender window is equal in size to receiver window
 - ▶ Starts at last byte read
- ▶ Sender Window:
 - ▶ Starts at last byte ACK'd
- ▶ Rules:
 - ▶ Window has fixed size of 3072 bytes
 - ▶ The first byte of the window is always the last ACK'd byte

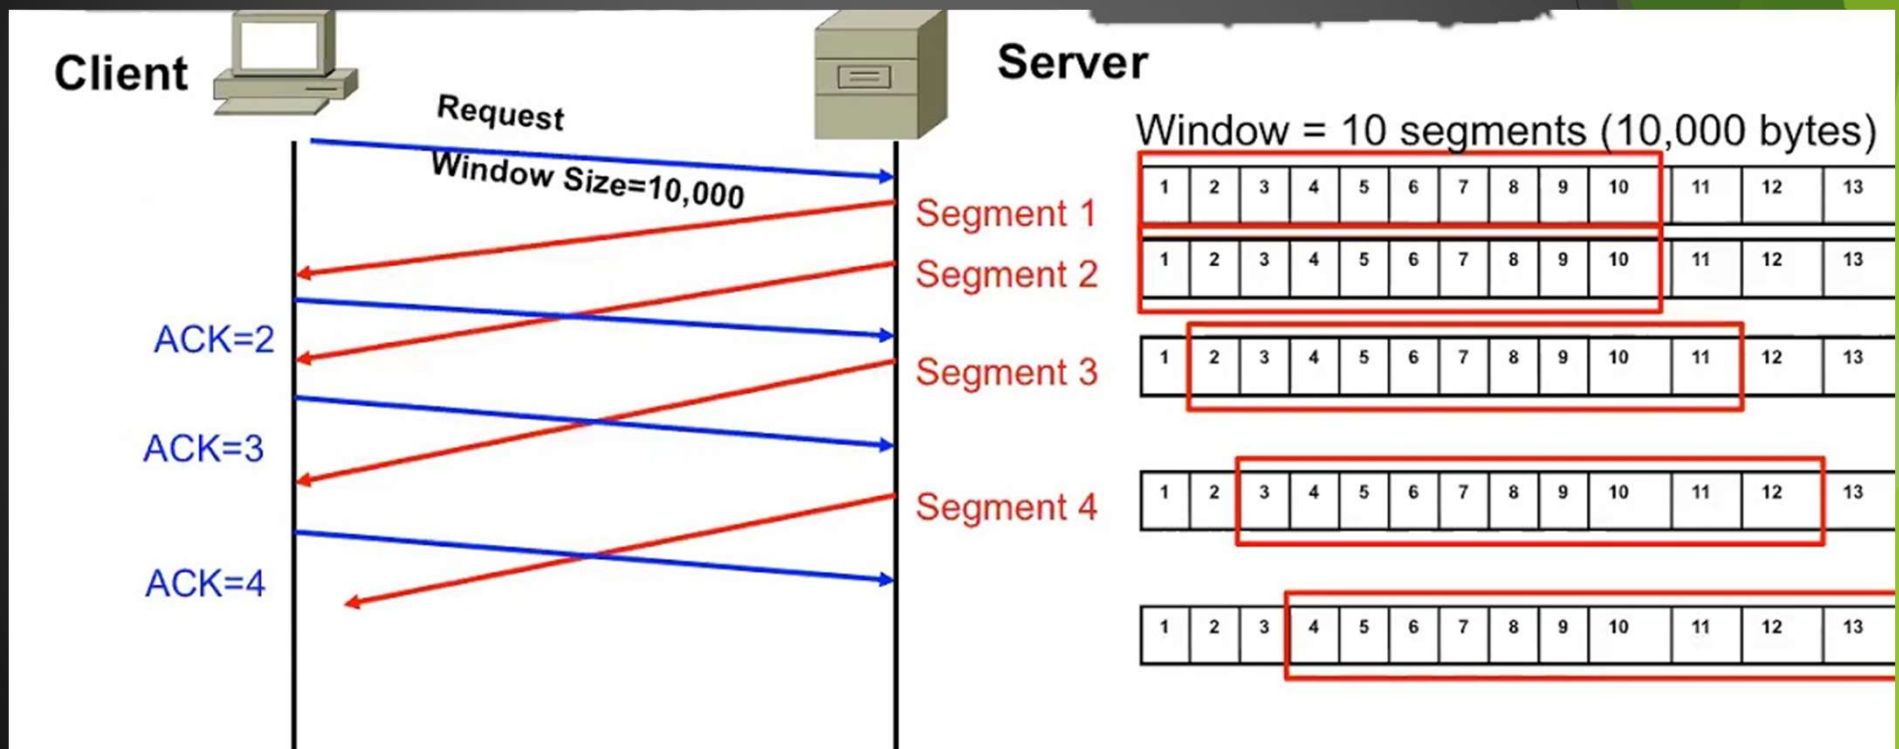
- `LastByteRead < NextByteExpected`
- `NextByteExpected <= LastByteRcvd + 1`
- Buffer bytes between `NextByteRead` and `LastByteRcvd`



- `LastByteAcked <= LastByteSent`
- `LastByteSent <= LastByteWritten`
- Buffer bytes between `LastByteAcked` and `LastByteWritten`



Sliding Window



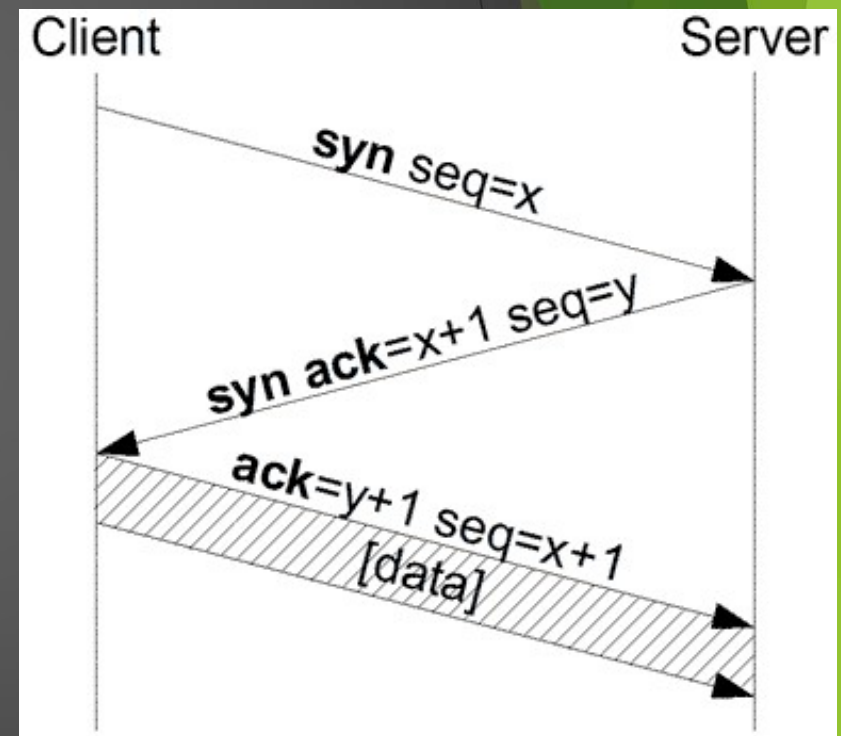
Advertised Window Size

- ▶ Receiver advertises how many bytes are left within its window
- ▶ Sender should not send packets outside of this window
- ▶ Advertised Window:
 - ▶ $\text{Size} = \text{Max Receiver Buffer} - ((\text{Next Expected Byte} - 1) - \text{Last Byte Read})$
 - ▶ Can assume that the $\text{Next Expected Byte} - 1 = \text{Last Byte Received}$
 - ▶ Since reliable connection
- ▶ Effective Window:
 - ▶ The sending window changes with the advertised window
 - ▶ $\text{Size} = \min(\text{Max Window Size}, \text{Advertised Window}) - (\text{Last Byte Sent} - \text{Last Byte ACK'd})$

Connection Handshake and Closing

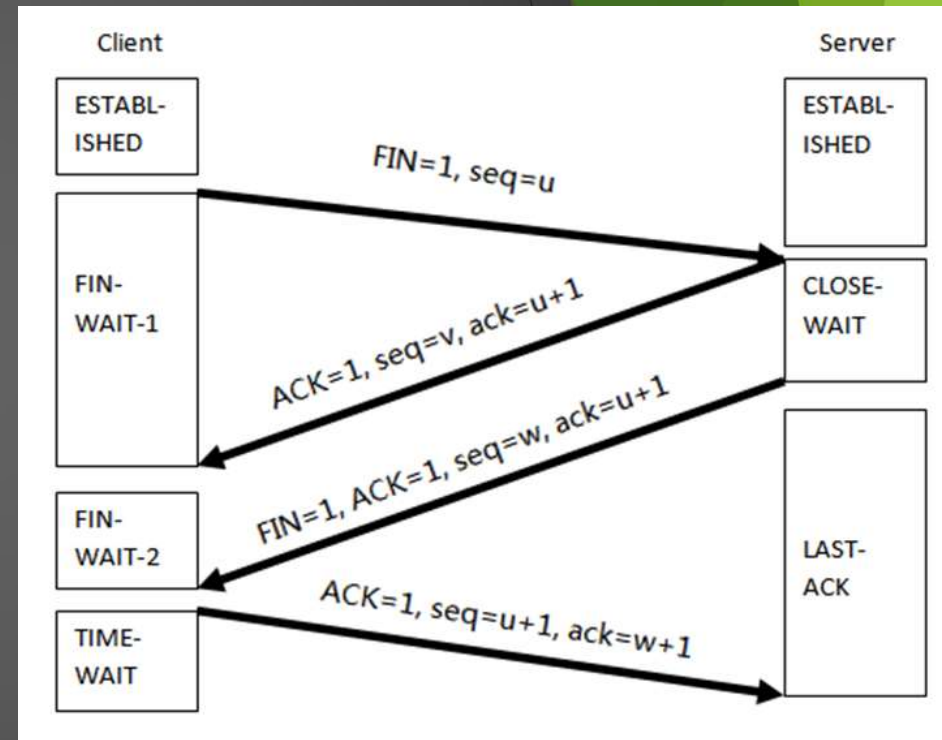
Network Initiation- 3 Way Handshake

- ▶ SYN:
 - ▶ Synchronize sequence numbers
 - ▶ Client gives server its initial sequence number
- ▶ SYN ACK:
 - ▶ Server replies with its own initial sequence number and the ACK'd client SYN
- ▶ ACK:
 - ▶ Client sends back server's ACK'd SYN



Network Termination- Teardown

- ▶ **FIN:**
 - ▶ Client sends packet with the flag set to indicate a teardown should begin
 - ▶ Packet may include data & receiver must be able to handle this
- ▶ **ACK, FIN:**
 - ▶ Server sends the ACK for the FIN
 - ▶ Sends its own FIN to indicate server teardown initiated
- ▶ **Final ACK:**
 - ▶ Client sends the ACK for the server FIN
 - ▶ Closes connection after sending
- ▶ **Server closes connection after receiving final ACK**



Getting Started

- ▶ Code to edit is in `transport.c`
 - ▶ See `transport.h` for function specifications
- ▶ Refer to `step_api.h` for functions used to communicate with network and application layers
 - ▶ Particularly `step_network_send()`, `step_network_recv()`, `step_app_recv()`, `step_app_send()`
- ▶ Test server and client are provided
 - ▶ Client sends path to a file with test input
 - ▶ Server opens file, reads input, and sends input back to client
 - ▶ Client prints what it received into file 'rcvd'
- ▶ Running:
 - ▶ Server:
 - ▶ `./server [port to listen to]`
 - ▶ Client:
 - ▶ `./client -f [file path] 127.0.0.1:[server port]`
 - ▶ Runs on Linux/Solaris. For other systems, use Vagrant
 - ▶ For newer MACs, use Docker as specified in the README